

Factorization in Deep Neural Networks - Part 2



IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom

Sessions

- 1 Deep Learning and Transfer Learning,
- 2 Quantification,
- 3 Pruning,
- 4 Factorization,
- 5 Fact. pt.2 : Operators and Architectures,
- 6 Distillation,
- 7 Embedded Software and Hardware for DL.
- 8 Presentations for challenge.

Sessions

- 1 Deep Learning and Transfer Learning,
- 2 Quantification,
- 3 Pruning,
- 4 Factorization,
- 5 **Fact. pt.2 : Operators and Architectures,**
- 6 Distillation,
- 7 Embedded Software and Hardware for DL.
- 8 Presentations for challenge.

Complexity of 2D Convolutions

$$N_{ops} = k.l.C_{in}C_{out}$$

with kernel size (k, l) , C_{in} input feature maps and C_{out} output feature maps.

To reduce the number of parameters, we can :

- Reduce the size of kernels
- Reduce the number of feature maps

Two strategies :

- Decompose kernels
- Depthwise convolutions

Complexity of 2D Convolutions

$$N_{ops} = k.l.C_{in}C_{out}$$

with kernel size (k, l) , C_{in} input feature maps and C_{out} output feature maps.

To reduce the number of parameters, we can :

- Reduce the size of kernels
- Reduce the number of feature maps

Two strategies :

- Decompose kernels
- Depthwise convolutions

Complexity of 2D Convolutions

$$N_{ops} = k.l.C_{in}C_{out}$$

with kernel size (k, l) , C_{in} input feature maps and C_{out} output feature maps.

To reduce the number of parameters, we can :

- Reduce the size of kernels
- Reduce the number of feature maps

Two strategies :

- Decompose kernels
- Depthwise convolutions

Complexity of 2D Convolutions

$$N_{ops} = k.l.C_{in}C_{out}$$

with kernel size (k, l) , C_{in} input feature maps and C_{out} output feature maps.

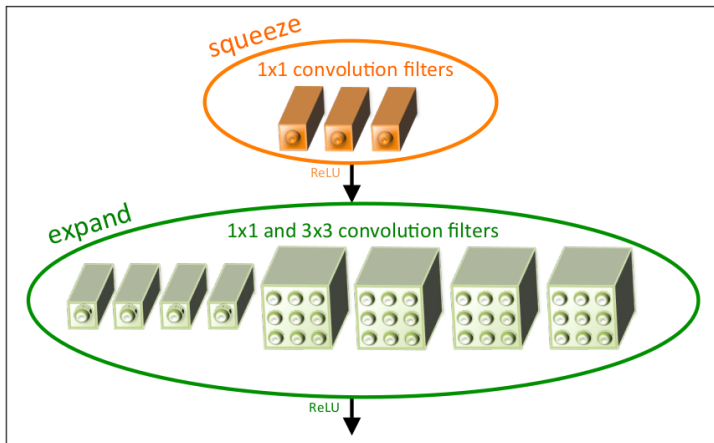
Decomposing kernels

Assuming $C_{in} = C_{out}$, decompose (k, l) kernel by $(k, 1)$ and $(1, l)$:

$$N_{ops} = k.1.C_{in}^2 + 1.l.C_{in}^2 = (l + k).C_{in}^2$$

with kernel size (k, l) , C_{in} input and out feature maps.

Introducing the Fire Module



landola et al. 2016, <https://arxiv.org/abs/1602.07360>

Depthwise convolutions

Instead of learning parameters that recombine all input feature maps to compute each feature maps, use "groups" of D input feature maps to compute $D C_{in}$ output feature maps.

Complexity of a Depthwise 2D Convolution

$$N_{ops} = k \cdot l \cdot C_{in} \cdot D \cdot C_{in}$$

with kernel size (k, l) , C_{in} feature maps and D is Depth.

Depthwise convolutions


Instead of learning parameters that recombine all input feature maps to compute each feature maps, use "groups" of D input feature maps to compute DC_{in} output feature maps.

Complexity of a Depthwise 2D Convolution


$$N_{ops} = k.l.C_{in}.D.C_{in}$$

with kernel size (k, l) , C_{in} feature maps and D is Depth.


MobileNetV1



MobileNetV2



Stride=1 block



Stride=2 block

<https://arxiv.org/abs/1704.04861> and <https://arxiv.org/abs/1801.04381>

Accuracy obtained on ImageNet

Network	Accuracy(%)	Params (M)
SqueezeNet	57.5	1.24
MobileNetV1	70.6	4.20
MobileNetV2	72.0	3.40

<https://arxiv.org/abs/1704.04861> and <https://arxiv.org/abs/1801.04381>

Alternatives to Convolution

Introducing Shift Attention Layer




Figure 1: Overview of the proposed method: we depict here the computation for a single output feature map d , considering a 1d convolution and its associated shift version. Panel (1) represents a standard convolutional operation: the weight filter $W_{d,c,\cdot}$ containing SC weights is moved along the spatial dimension (L) of the input to produce each output in Y_d . In panel (2), we depict the attention tensor A on top of the weight filter: the darker the cell, the most important the corresponding weight has been identified to be. At the end of the training process, A should contain only binary values with a single 1 per slice $A_{d,c,\cdot}$. In panel (3), we depict the corresponding obtained shift layer: for each slice along the input feature maps (C), the cell with the highest attention is kept and the others are disregarded. As a consequence, the initial convolution with a kernel size S has been replaced by a convolution with a kernel size 1 on a shifted version of the input X . As such, the resulting operation in panel (3) is exactly the same as the shift layer introduced in Wu et al. [2017], but here the shifts have been trained instead of being arbitrarily predetermined.

Alternatives to Convolution




Figure 7: Evolution of accuracy when applying compression methods on different DNN architectures trained on CIFAR10.

Hacene et al. 2019, <https://arxiv.org/abs/1905.12300>